

# Weighted Finite State Transducers in Automatic Speech Recognition

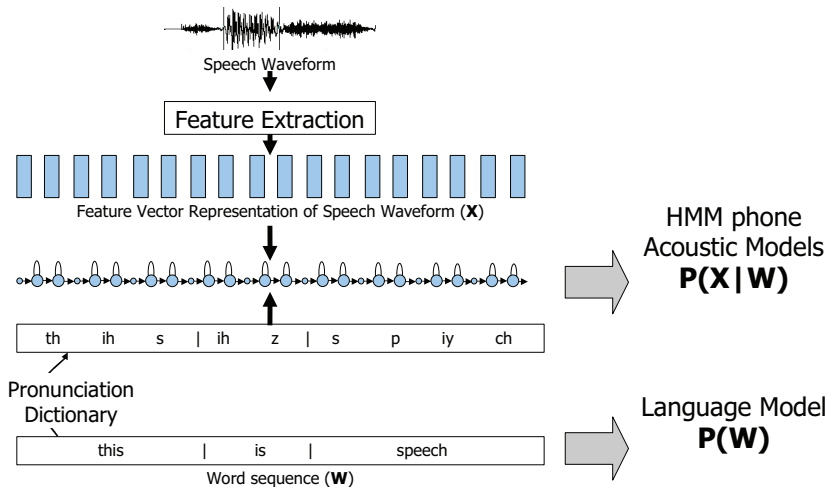
ZRE lecture 10.04.2013

Mirko Hannemann

Slides provided with permission, Daniel Povey  
some slides from T. Schultz, M. Mohri and M. Riley

10.04.2013

# Automatic speech recognition




# Automatic speech recognition

- Determine the most probable word sequence  $\tilde{W}$  given the observed acoustic signal  $Y$


$$\tilde{W} = \operatorname{argmax} P(W|Y) = \frac{\operatorname{argmax} P(W) P(Y|W)}{P(Y)}$$

- Search for word sequence  $\tilde{W}$  that maximizes  $P(W)$  and  $P(Y|W)$ 
  - $P(W)$  = language model (likelihood of word sequence)
  - $P(Y|W)$  = acoustic model  
(likelihood of observed acoustic signal given word sequence)


# Hidden Markov Model, Token passing



# Viterbi algorithm, trellis



# Decoding graph/recognition network




# Why Finite State Transducers?

## Motivation:

- most components (LM, lexicon, lattice) are finite-state
  - unified framework for describing models
  - integrate different models into a single model via composition operations
  - improve search efficiency via optimization algorithms
  - flexibility to extend (add new models)
- *speed*: pre-compiled search space, near realtime performance on embedded systems
- *flexibility*: same decoder used for hand-held devices and LVCSR


# Finite State Acceptor (FSA)



- ▶ An FSA “accepts” a set of strings
- ▶ (a string is a sequence of symbols).
- ▶ View FSA as a representation of a possibly infinite set of strings.
- ▶ This FSA accepts just the string  $ab$ , i.e. the set  $\{ab\}$
- ▶ Numbers in circles are state labels (not really important).
- ▶ Labels on arcs are the symbols.
- ▶ Start state(s) bold; final/accepting states have extra circle.
  - ▶ Note: it is sometimes assumed there is just one start state.




# A less trivial FSA




- ▶ The previous example doesn't show the power of FSAs because we could represent the set of strings finitely.
- ▶ This example represents the infinite set  $\{ab, aab, aaab, \dots\}$
- ▶ Note: a string is “accepted” (included in the set) if:
  - ▶ There is a path with that sequence of symbols on it.
  - ▶ That path is “successful” (starts at an initial state, ends at a final state).

# The epsilon symbol



- ▶ The symbol  $\epsilon$  has a special meaning in FSAs (and FSTs)
- ▶ It means “no symbol is there”.
- ▶ This example represents the set of strings  $\{a, aa, aaa, \dots\}$
- ▶ If  $\epsilon$  were treated as a normal symbol, this would be  $\{a, a\epsilon a, a\epsilon a\epsilon a, \dots\}$
- ▶ In text form,  $\epsilon$  is sometimes written as  $\langle \text{eps} \rangle$
- ▶ Toolkits implementing FSAs/FSTs generally assume  $\langle \text{eps} \rangle$  is the symbol numbered zero

# Weighted finite state acceptors



- ▶ Like a normal FSA but with costs on the arcs and final-states
- ▶ Note: cost comes after “/”. For final-state, “2/1” means final-cost 1 on state 2.
- ▶ View WFSA as a function from a string to a cost.
- ▶ In this view, unweighted FSA is  $f : \text{string} \rightarrow \{0, \infty\}$ .
- ▶ If multiple paths have the same string, take the one with the lowest cost.
- ▶ This example maps  $ab$  to  $(3 = 1 + 1 + 1)$ , all else to  $\infty$ .

# Semirings


- ▶ The semiring concept makes WFSAs more general.
- ▶ A semiring is
  - ▶ A set of elements (e.g.  $\mathbb{R}$ )
  - ▶ Two special elements  $\bar{1}$  and  $\bar{0}$  (the identity element and zero)
  - ▶ Two operations,  $\oplus$  (plus) and  $\otimes$  (times) satisfying certain axioms.

*Semiring examples.*  $\oplus_{\log}$  is defined by:  $x \oplus_{\log} y = -\log(e^{-x} + e^{-y})$ .

SEMIRING	SET	$\oplus$	$\otimes$	$\bar{0}$	$\bar{1}$
Boolean	$\{0, 1\}$	$\vee$	$\wedge$	0	1
Probability	$\mathbb{R}_+$	+	$\times$	0	1
Log	$\mathbb{R} \cup \{-\infty, +\infty\}$	$\oplus_{\log}$	+	$+\infty$	0
Tropical	$\mathbb{R} \cup \{-\infty, +\infty\}$	min	+	$+\infty$	0


- In WFSAs, weights are multiplied along paths
- summed over paths with identical symbol-sequences

# Weights vs. costs




- ▶ Personally I use “cost” to refer to the numeric value, and “weight” when speaking abstractly, e.g.:
  - ▶ The acceptor above accepts  $a$  with unit weight.
  - ▶ It accepts  $a$  with zero cost.
  - ▶ It accepts  $bc$  with cost  $4 = 2 + 1 + 1$
  - ▶ State 1 is final with unit weight.
  - ▶ The acceptor assigns zero weight to  $xyz$ .
  - ▶ It assigns infinite cost to  $xyz$ .

# Weights and costs




- ▶ Consider the WFSA above, with the tropical (“Viterbi-like”) semiring. Take the string *ab*.
- ▶ We “multiply” ( $\otimes$ ) the weights along paths; this means adding the costs.
- ▶ Two paths for *ab*:
  - ▶ One goes through states (0, 1, 1); cost is  $(0 + 1 + 2) = 3$
  - ▶ One goes through states (0, 2, 3); cost is  $(1 + 2 + 2) = 5$
- ▶ We add weights across different paths; tropical  $\oplus$  is “take min cost” → this WFSA maps *ab* to 3

# Weighted finite state transducers (WFST)



- ▶ Like a WFSA except with two labels on each arc.
- ▶ View it as a function from a (pair of strings) to a weight
- ▶ This one maps  $(a, x)$  to 1 and all else to  $\infty$
- ▶ Note: view 1 and  $\infty$  as costs.  $\infty$  is  $\bar{0}$  in semiring.
- ▶ Symbols on the left and right are termed “input” and “output” symbols.

# Composition of WFSTs



- ▶ Notation:  $C = A \circ B$  means,  $C$  is  $A$  composed with  $B$ .
- ▶ In special cases, composition is similar to function composition
- ▶ Composition algorithm “matches up” the “inner symbols”
  - ▶ i.e. those on the output (right) of  $A$  and input (left) of  $B$



# Composition algorithm

- ▶ Ignoring  $\epsilon$  symbols, algorithm is quite simple.
- ▶ States in  $C$  correspond to tuples of (state in  $A$ , state in  $B$ ).
  - ▶ But some of these may be inaccessible and pruned away.
- ▶ Maintain queue of pairs, initially the single pair  $(0, 0)$  (start states).
- ▶ When processing a pair  $(s, t)$ :
  - ▶ Consider each pair of (arc  $a$  from  $s$ ), (arc  $b$  from  $t$ ).
  - ▶ If these have matching symbols (output of  $a$ , input of  $b$ ):
    - ▶ Create transition to state in  $C$  corresponding to (next-state of  $a$ , next-state of  $b$ )
    - ▶ If not seen before, add this pair to queue.
- ▶ With  $\epsilon$  involved, need to be careful to avoid redundant paths...

# Construction of decoding network

- WFST approach [Mohri et al.]
- exploit several knowledge sources (lexicon, grammar, phonetics) to find most likely spoken word sequence

$$HCLG = H \circ C \circ L \circ G \quad (1)$$

- G** probabilistic grammar or language model acceptor (word)
- L** lexicon (phones to words)
- C** context-dependent relabeling (ctx-dep-phone to phone)
- H** HMM structure (PDF labels to context-dependent phones)

Create H, C, L, G separately and compose them together

# Language model

Estimate probability of a word sequence  $W$ :

$$P(W) = P(w_1, w_2, \dots, w_N) \quad (2)$$

$$P(W) = P(w_1) \cdot P(w_2|w_1) \cdot P(w_3|w_1, w_2) \cdot \dots \cdot P(w_N|w_1, \dots, w_{N-1}) \quad (3)$$

Approximate by sharing histories  $h$  (e.g. bigram history):

$$P(W) \approx P(w_1) \cdot P(w_2|w_1) \cdot P(w_3|w_2) \cdot \dots \cdot P(w_N|w_{N-1}) \quad (4)$$

What to do if a certain history was not observed in training texts?


- statistical smoothing of the distribution
- interpolation of different orders of history
- backing-off to shorter history

# Language model acceptor G

- **G:** Grammar Transducer  
 Backing-off language model:

$$p(w|h) = \begin{cases} f(w|h) & : \text{if } N(w, h) > 0 \\ \alpha(h) \cdot f(w|\bar{h}) & : \text{if } N(w, h) = 0 \end{cases}$$

- **Input:** word
- **Weight:** history dependent word probability




# Language models (ARPA back-off)

\1-grams:

-5.2347 a -3.3  
-3.4568 b  
0.0000 <s> -2.5  
-4.3333 </s>

\2-grams:

-1.4568 a b  
-1.3049 <s> a  
-1.78 b a  
-2.30 b </s>




# Pronunciation lexicon L

A	ax #1
ABERDEEN	ae b er d iy n
ABOARD	ax b r ao dd
ADD	ae dd #1
ABOVE	ax b ah v

Added disambiguation symbols:



- if a phone sequences can output different words ("I scream for ice cream.")
- non-determinism: introduce disambiguation symbols, remove at last stage

# Deterministic WFSTs





- ▶ Taken to mean “deterministic on the input symbol”
- ▶ I.e., no state can have  $> 1$  arc out of it with the same input symbol.
- ▶ Some interpretations (e.g. Mohri/AT&T/OpenFst) allow  $\epsilon$  input symbols (i.e. being  $\epsilon$ -free is a separate issue).
- ▶ I prefer a definition that disallows epsilons, except as necessary to encode a string of output symbols on an arc.
- ▶ Regardless of definition, not all WFSTs can be determinized.

# Determinization (like making tree-structured lexicon)







# Minimal deterministic WFSTs




- ▶ Here, the left FSA is not minimal but the right one is.
- ▶ “Minimal” is normally only applied to *deterministic* FSAs.
- ▶ Think of it as suffix sharing, or combining redundant states.
- ▶ It's useful to save space (but not as crucial as determinization, for ASR).

# Minimization (like suffix sharing)



# Pronunciation lexicon L


- **L:** Context-Dependency Transducer
  - **Input:** context-independent phone (phoneme)
  - **Output:** word
  - **Weight:** pronunciation probability



# Context-dependency of phonemes

So far we use phonemes independent of context, but:

- co-articulation: pronunciation of phones changes with surrounding phones



# Context-dependency transducer C


Introduce context-dependent phones:

- tri-phones: each model depends on predecessor and successor phoneme:  $a-b-c$  ( $b/a_c$ )
- implemented as context-dependency transducer C

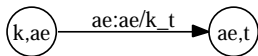
**Input:** context-dependent phone (triphone)

**Output:** context-independent phone (phone)


- shown is one path of it



## Context-dependency transducer C

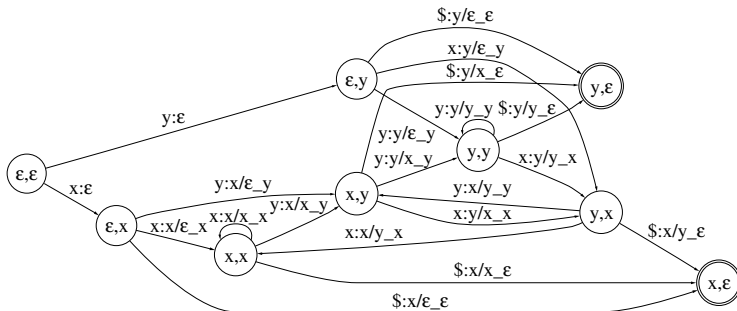


(a)




(b)

Figure 6: Context-dependent triphone transducer transition: (a) non-deterministic, (b) deterministic.




# HMM as transducer



# HMM as transducer (monophone)

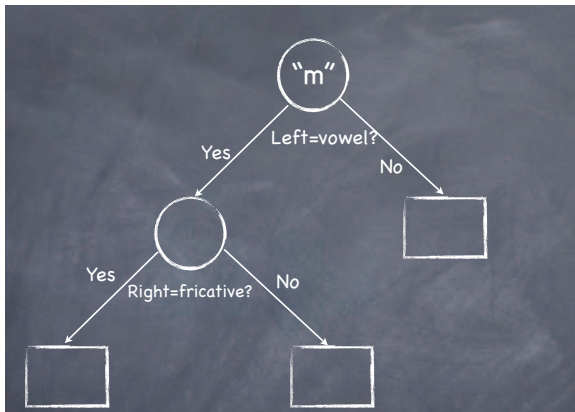
- H: HMM Topology Transducer (maps states to phonemes)
  - **Input:** state
  - **Output:** context-dependent phone (triphone)
  - **Weight:** HMM transition probability



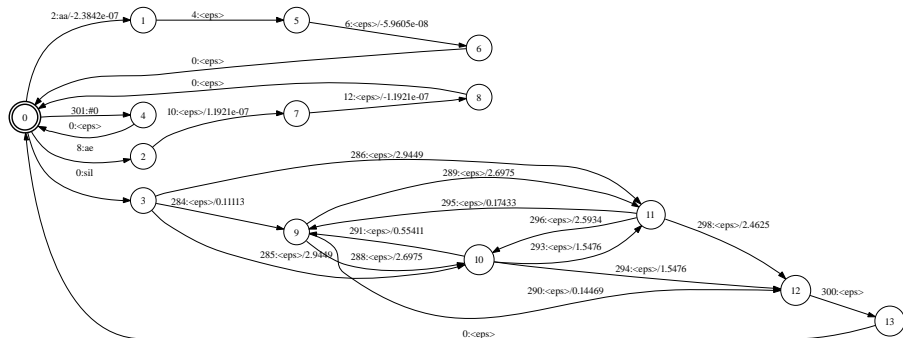


# Phonetic decision tree

- too many context-dependent models ( $N^3$ )  $\rightarrow$  clustering
- determine model-id (gaussian) based on phoneme context and state in HMM
- using questions about context (sets of phones)



# HMM transducer $H_a$



(here shown for monophone case, without self-loops)

# Construction of decoding network

WFST approach by [Mohri et al.]

$$HCLG = rds(min(det(H \circ det(C \circ det(L \circ G)))))) \quad (5)$$

rds - remove disambiguation symbols

min - minimization, includes weight pushing

det - determinization

Kaldi toolkit [Povey et al.]

$$HCLG = asl(min(rds(det(H_a \circ min(det(C \circ min(det(L \circ G)))))))) \quad (6)$$

asl - add self loops

rds - remove disambiguation symbols

# Decoding graph construction (complexities)

- Have to do things in a careful order or algorithms "blow up"
- Determinization for WFSTs can fail
  - need to insert "disambiguation symbols" into the lexicon.
  - need to "propagate these through" H and C.
- Need to guarantee that final HCLG is stochastic:
  - i.e. sums to one, like a properly normalized HMM
  - needed for optimal pruning (discard unlikely paths)
  - usually done by weight-pushing, but standard algorithm can fail, because FST representation of back-off LMs is non-stochastic
- We want to recover the phone sequence from the recognized path (words)
  - sometimes also the model-indices (PDF-ids) and the HCLG arcs that were used in best path

# Decoding with WFSTs (finding best path)

- Solve


$$W' = \operatorname{argmax}_W P(X|W) P(W)$$

- Compose recognizer as  $(H \circ C \circ L \circ G)$  which maps states to word sequences
- Decode by aligning the feature vectors  $X$  with HCLG "

i.e.,


$$W' = \operatorname{argmax}_W X \circ (H \circ C \circ L \circ G)$$

# Decoding with WFSTs



- ▶ First– a “WFST definition” of the decoding problem.
- ▶ Let  $U$  be an FST that encodes the acoustic scores of an utterance (as above).
- ▶ Let  $S = U \circ HCLG$  be called the search graph for an utterance.
- ▶ Note: if  $U$  has  $N$  frames (3, above), then
  - ▶ #states in  $S$  is  $\leq (N + 1)$  times #states in  $HCLG$ .
  - ▶ Like  $N + 1$  copies of  $HCLG$ .

# Viterbi algorithm, trellis



# Decoding with WFSTs

- ▶ With beam pruning, we search a subgraph of  $S$ .
- ▶ The set of “active states” on all frames, with arcs linking them as appropriate, is a subgraph of  $S$ .
- ▶ Let this be called the *beam-pruned subgraph* of  $S$ ; call it  $B$ .
- ▶ A standard speech recognition decoder finds the best path through  $B$ .
- ▶ In our case, the output of the decoder is a linear WFST that consists of this best path.
- ▶ This contains the following useful information:
  - ▶ The word sequence, as output symbols.
  - ▶ The state alignment, as input symbols.
  - ▶ The cost of the path, as the total weight.




# Decoding output

```
utt1 [ 2 6 6 6 6 10 ] [ 614 613 613 613 711 ] [ 122 123
utt1 SIL th ax
utt1 <s> THE
```

# Word Lattice / Word Graph

Word Lattice: a compact representation of the search space



# Lattices as WFSTs

The word "lattice" is used in the ASR literature as:

- Some kind of compact representation of the alternate word hypotheses for an utterance.
- Like an N-best list but with less redundancy.
- Usually has time information, sometimes state or word alignment information.
- Generally a directed acyclic graph with only one start node and only one end node.

# Lattice Generation with WFSTs [Povey12]

Basic process (not memory efficient):

- Generate beam-pruned subgraph B of search graph S
  - The states in B correspond to the active states on particular frames.
  - Prune B with beam  $\alpha$  to get pruned version P.
  - Convert P to acceptor and lattice-determinize to get A (deterministic acceptor)
- No two paths in L have same word sequence (take best)
- Prune A with beam  $\alpha$  to get final lattice L (in acceptor form).

# Finite State Transducers for ASR

## Pro's:

**Fast:** compact/minimal search space due to combined minimization of lexicon, phonemes, HMM's

**Simple:** easy construction of recognizer by composition from states, HMMs, phonemes, lexicon, grammar

**Flexible:** whatever new knowledge sources, the compose/optimize/search remains the same

## Con's:

- composition of complex models generates a huge WFST
- search space increases, and huge memory is required
- esp. how to deal with huge language models

# Ressources

- OpenFST** <http://www.openfst.org/> Library, developed at Google Research (M. Riley, J. Schalkwyk, W. Skut) and NYU's Courant Institute (C. Allauzen, M. Mohri)
- Mohri08** M. Mohri et al., *"Speech Recognition with weighted finite state transducers."*
- Povey11** D. Povey et al., *"The Kaldi Speech Recognition Toolkit."*
- Povey12** D. Povey et al., *"Generating exact lattices in the WFST framework."*